


GIC Certificate of Model Health

Independent AI Model Health Assessment · Ed25519 Signed · Independently Verifiable
Certificate ID: GIC-PRIME56-2026-001 · Issued: 2026-08-19 · Valid: perpetual (point-in-time measurement)

MODEL IDENTITY

Certified Name	Prime-56
Base Model	Qwen3-32B (Apache 2.0) by Alibaba
Optimization	Conductor-mapped surgical optimization by Growing Intelligence
Parameters	32 billion, bf16, 64 layers, 14 shards
Base Weights SHA256	73ef8f2a7629f5e28f132eab83fcc77ee1230f74...
Certified Weights SHA256	9359a559c7630004807e3debe86dc1e325109ede0a2eaf270f8f9ce7e994e53e
Conductors	2,377,576 live of 2,490,368 (95.5%)
Scan Questions	10,000 (2,500 × D1/D3/D4/D5)
Evidence Questions	GEN56 (D2/D3/D4) + SAFE56 (D5) — contamination-proof

MODEL MRI — CONDUCTOR FIELD VISUALIZATION



Conductor field — 2,377,576 live conductors reading the model across 64 layers. Each cell represents a sampled conductor position. Color indicates verdict at that moment during the scan.

Correct (top-1) Fixable (top-2 to top-20) Incorrect (beyond top-20) Active read

RESULTS BY DIMENSION — BEFORE & AFTER OPTIMIZATION

Dimension	N	Correct Before	Correct After	Δ	Fixable	Repaired
D1 Knowledge	2,500	7.4%	22.6%	+380	538	380
D3 Skills	2,500	31.2%	39.9%	+217	585	217
D4 Logic	2,500	3.6%	8.6%	+124	244	124
D5 Safety	2,500	11.1%	15.6%	+111	520	111
D2 Reasoning	DEGRADED — measured under single-token constraint (see disclosure)					

56-ELEMENT FUNCTIONAL HEALTH PROFILE

Each element represents the intersection of one atom (row) and one law (column) from the Holosynthics framework — 8 atoms × 7 laws = 56 functional elements. Colors indicate health status after optimization.

Atom ↓ · Law →	Distinction	Feedback	Criticality	Composition	Complexity	Network	Uncertainty
Bool	24%	18%	0%	N/A	33%	50%	—
Number	34%	10%	19%	22%	33%	0%	N/A
Symbol	27%	12%	13%	11%	12%	10%	—
Identity	24%	6%	8%	N/A	—	29%	—
Relation	21%	9%	12%	2%	N/A	19%	—
Time	24%	4%	1%	6%	6%	0%	N/A
Space	28%	6%	—	1%	—	N/A	—
Entropy	11%	14%	33%	—	—	—	—

Green = healthy (>70% correct after optimization). Amber = partially optimized (30-70%). Red = weak (<30%). Gray = not measured in this build. Percentages populated from scan results.

ELEMENT RANKING — TOP 10 BY OBSERVATION COUNT

#	Element	Observations	Correct %	Fixable %
1	Time:Criticality	1,768	0.8%	5.4%
2	Symbol:Feedback	1,105	12.3%	18.8%
3	Symbol:Distinction	1,003	27.4%	15.1%
4	Time:Feedback	617	3.6%	20.9%
5	Time:Distinction	605	23.7%	29.0%
6	Number:Distinction	344	34.5%	34.5%
7	Number:Complexity	208	32.7%	37.9%
8	Space:Composition	182	0.8%	22.7%
9	Identity:Distinction	152	24.0%	30.1%
10	Number:Feedback	124	10.2%	26.1%

OPTIMIZATION SUMMARY

Metric	Value	Notes
Total scanned	10,000	10,000 questions, 4 dimensions
Fixable found (top-20)	1,887 (18.9%)	Top-5: 870 · Top 6→20: 1,138
Successfully repaired	832 / 1,837 attempted (45.3%)	Promoted to top-1
Net correct gained	+832	Before: 1,334 → After: 2,166 (rank-measured promotions; reconstruction fidelity disclosed)
Collateral damage	0	0 events across 1,837 gated installs · control set grew 1,334 → 2,166

Zero collateral damage — verified. 2,166 control items tested (1,334 scan-correct + 832 kept flips). Zero previously-correct answers changed by the optimization. Wilson 95% confidence interval: 0 events / 1,837 gated installs. Binary-preference items excluded from the control set — these are the only item class that has ever exhibited collateral under minimal edits. This exclusion is a measurement choice and is stated here for transparency.

EVIDENCE INTEGRITY

Contamination-Proof Evidence (GEN56 / SAFE56)

Source	Dimension	Questions	Before	After	Δ
GEN56 Conservation	D3	1,000	79.2%	80.2%	+10
GEN56 Control	D4	1,000	0.0%	0.1%	+1
SAFE56	D5	1,000	65.1%	65.6%	+5
GEN56 Entropy	D2	DEGRADED — see disclosure			

Statistical Tests

Test	Result
ConStat (contamination)	structural: seed commitments published before generation + per-row canaries (ConStat not run — no implementation in the pipeline; disclosed)
McNemar (paired before/after)	n.s. per dimension (D3 p=0.54, D2 p=0.50, D4 p=1.00, D5 p=0.66) — no degradation; scan-set gains carried by 832 rank-measured promotions
Wilson CI (collateral)	0 events / 1,837 installs — consistent with true zero

DISCLOSURES

D2 Reasoning — DEGRADED: Measured under single-token constraint; effective measurement covers the extractable digit, not the full answer. GEN56 produces multi-token fraction golds (e.g., 1007/23742). Scoring the first token is a partial measurement, stated honestly.

Binary exclusion: Binary-preference (reachable-type) items excluded from the control set. These are the only item class that has ever exhibited collateral under minimal edits (magnitude 0.96 against a promoted mean of 80.3). The "zero collateral" claim is measured on the remaining non-binary control items.

Unscoreable budget: 11.4% of questions exceeded the 2,048-token reasoning budget and are excluded from the denominator. The effective sample is 8,863 of 10,000 attempted.

This document records a named reproducible procedure and its measured results. It is not an accreditation and invokes no external standards body.

CRYPTOGRAPHIC PROVENANCE — HASH CHAIN

The chain below proves the certified weights are derived from the frozen base through a documented repair process, without revealing the repair method.

BASE WEIGHTS (frozen)
 sha256: 73ef8f2a7629f5e28f132eab83fcc77ee1230f74ed2c910acd6ace553399426f
 ————— fixes.pkl —————→

REPAIR RECORD
 sha256: ee585d9d12ba0340b7e1ba80621e4140b6f7736f758b88cfd55e68a00da2dd2e
 ————— merge + strip —————→

CERTIFIED WEIGHTS (Prime-56)
 sha256: 9359a559c7630004807e3debe86dc1e325109ede0a2eaf270f8f9ce7e994e53e

CONDUCTOR MAP (referenced by hash, never published)
 sha256: 87a59661112914a2fc...

MERKLE ROOT (over all evidence triples)
 sha256: f17f61c159cac46d14bf918063f31413e7c269bf6402a1b277f6c18c66515f88

EVIDENCE TRAIL

Artifact	SHA256 (first 16)	Location
Freeze declaration	73ef8f2a7629f5e2	V1 Railway · V2 GitHub · V3 Arweave
Conductor map (.npz)	87a59661112914a2	Private — never published
Entity (14 shards)	9359a559c7630004	Private — never published
fixes.pkl	ee585d9d12ba0340	Private — never published
Scan results (parquet)	f9f057a29d3433ad	Private
GEN56 seed commitment	d3d72e7cb19eea10	Published before generation
SAFE56 seed commitment	2bf50dcc292c5669	Published before generation
Certificate (this document)	91fae5d41158b304	Public — proof-vault

Ed25519 Digital Signature

SIGNATURE

7e3049f4e39b922211b6dd85087e8bbde6d348c7a1b207a93ddc24e7d1c130b1714fb80bfc7e09384565e3d4de4fa4bb6b561cd55c76c97efe0d063382a3107

PUBLIC KEY (Founder)

FtrWshUc/9rg5Cz+ARi5DP/yyqFhWMJLmx0VHKc3wpk=

MERKLE ROOT

f17f61c159cac46d14bf918063f31413e7c269bf6402a1b277f6c18c66515f88



→
proof-
vault

Verify: download the certificate JSON from the proof-vault, compute the Merkle root, and verify the Ed25519 signature against the public key above. If the signature is valid, the certificate is authentic and unmodified.

Growing Intelligence · jonathan@growing-intelligence.com · growing-intelligence.com

This document records a named reproducible procedure and its measured results. It is not an accreditation and invokes no external standards body.