



Growing Intelligence

GIC Certificate of Model Health

Independent AI Model Health Assessment · Ed25519 Signed · Independently Verifiable

Certificate ID: GIC-PRIMEBANK-2026-001 · Issued: 2026-08-24 · Valid: perpetual (point-in-time measurement)

MODEL IDENTITY

Certified Name	Prime-Bank
Base Model	DictaLM-3.0-24B-Thinking by Dicta
Optimization	Conductor-mapped surgical optimization by Growing Intelligence
Parameters	23.57 billion, bfloat16, 40 layers
Conductors	204,800 (5,120 hidden × 40 layers)
Scan Questions	1,779 (HebrewBankBench v1)
Evidence Questions	held-out Hebrew banking + general — scored on first-token rank over the full vocabulary

MODEL MRI — CONDUCTOR FIELD VISUALIZATION



Conductor field — 204,800 conductors reading the model across 40 layers, layer 30 repair band highlighted.

RESULTS BY DIMENSION — BEFORE & AFTER OPTIMIZATION

Dimension	N	Correct Before	Correct After	Δ
D3	546	33.0%	72.9%	+218
D1	279	38.7%	74.6%	+100
D5	212	31.1%	67.5%	+77
D2	199	41.7%	69.8%	+56
D4	168	23.2%	50.0%	+45
D6	125	44.8%	76.8%	+40
D7	125	0.8%	69.6%	+86
D8	125	25.6%	66.4%	+51



Growing Intelligence

56-ELEMENT FUNCTIONAL HEALTH PROFILE

Each element is the intersection of one atom (row) and one law (column) — 8 atoms × 7 laws = 56 functional elements; this corpus observes 35 of them, 444 of 1779 questions (25.0%) unmapped by the classifier. Cell shows correct% AFTER optimization.

Atom ↓ · Law →	Identity	Conserve	Reduce	Compose	Symmetry	Transform	Converge
Bool	0.0%	80.5%	81.3%	—	—	—	—
Num	55.6%	50.0%	33.3%	69.2%	—	87.5%	100.0%
Sym	92.3%	100.0%	—	100.0%	—	75.0%	78.9%
Rel	84.8%	—	—	—	—	—	—
ID	75.5%	76.9%	—	50.0%	100.0%	—	100.0%
Ent	68.8%	—	58.3%	66.7%	94.3%	83.3%	—
Time	63.9%	66.7%	46.9%	0.0%	—	—	100.0%
Space	68.9%	—	57.1%	76.9%	0.0%	—	74.3%

Green >70% correct after. Amber 30-70%. Red <30%. Gray = not observed by this corpus.

ELEMENT RANKING — TOP 20 BY OBSERVATION COUNT

#	Element	Observations	Correct% Before	Correct% After
1	Ent.Identity	346	39.3%	68.8%
2	Num.Identity	216	27.3%	55.6%
3	Bool.Conserve	169	32.0%	80.5%
4	Bool.Reduce	107	55.1%	81.3%
5	Time.Reduce	49	20.4%	46.9%
6	ID.Identity	49	46.9%	75.5%
7	Rel.Identity	46	37.0%	84.8%
8	Space.Identity	45	22.2%	68.9%
9	Time.Identity	36	44.4%	63.9%
10	Space.Converge	35	37.1%	74.3%
11	Ent.Symmetry	35	54.3%	94.3%
12	ID.Conserve	26	15.4%	76.9%
13	Num.Compose	26	23.1%	69.2%
14	Num.Reduce	24	12.5%	33.3%
15	Sym.Converge	19	10.5%	78.9%
16	Ent.Transform	18	22.2%	83.3%
17	Time.Conserve	15	6.7%	66.7%
18	Sym.Identity	13	30.8%	92.3%
19	Space.Compose	13	38.5%	76.9%
20	Ent.Reduce	12	16.7%	58.3%



Growing Intelligence

OPTIMIZATION SUMMARY

Metric	Value	Notes
Total scanned	1,779	1,779 questions, 8 dimensions
Fixable found (rank 2+)	1,214	reclassified after float32-corrected scan (8 of the original 573 "correct" were misclassified)
Successfully repaired	673 of 1,214 fixable (55.4%)	Promoted to rank 1
Net correct gained	+673	Before: 565 → After: 1238
Collateral damage	0	0 above a float32-verified floor of 565/573

Zero collateral above the harness floor — verified. All 573 originally-scanned controls were re-scored in float32 with the corrected batched-regime ranking. 8 of them do not reproduce rank 1 on the untouched model — an original-scan misclassification (not repair-caused damage) — so that is the measurement floor. The remaining 565 controls all hold at rank 1 with all 1,989 repair vectors installed.

EVIDENCE INTEGRITY — PER-DIMENSION, SAME QUESTIONS, SAME PROMPT, BEFORE AND AFTER

Dimension	Questions	Before	After	Δ
D3	546	33.0%	72.9%	+218
D1	279	38.7%	74.6%	+100
D5	212	31.1%	67.5%	+77
D2	199	41.7%	69.8%	+56
D4	168	23.2%	50.0%	+45
D6	125	44.8%	76.8%	+40
D7	125	0.8%	69.6%	+86
D8	125	25.6%	66.4%	+51

STATISTICAL TESTS

Test	Result
Collateral measurement	Census, not a sample — all 573 originally-scanned controls re-scored with the final 1,989 vectors installed, in float32.
Harness floor	8 of 573 fail to reproduce rank 1 on the untouched model, measured before any edit and excluded from the control set every collateral figure is measured against (565/573 remain as the verified floor).
Monitoring neutrality	ON == OFF verified bit-identical on the calibration set: maximum absolute logit difference 0.0, argmax identical 24 of 24.
Measurement bugs	Two ranking bugs found and fixed before this certificate's numbers were measured (see disclosures). A residual bf16 exact-tie effect was also found; this certificate is measured in float32.
Significance testing	Not run. The reported change is a direct count over the full question set, not an inference from a sample.

DISCLOSURES

- **Collateral is measured above a float32-verified floor.** 8 of the original 573 controls were reclassified from "correct" to "fixable" after corrected measurement showed they never reached rank 1 — an original-scan artefact, not repair-caused damage.
- **v2's 799 vectors are served as a runtime plugin, not baked,** because their gate (a 56-dim fingerprint cosine) has no representation as the scalar-projection technique used to bake v1.
- **The attack pre-filter is not a financial-advice filter.** Its 72 rules detect prompt injection and jailbreaks — verified 0.0% false positives on held-out testing.
- **Accuracy definition.** A question is correct when the gold answer reaches rank 1 over the full 131,072-token vocabulary at the answer position, under a fixed brief-answer instruction applied identically before and after.
- **Method.** The optimization procedure is patent-protected and is not described in this document.



Growing Intelligence

CRYPTOGRAPHIC PROVENANCE — HASH CHAIN

The chain below proves the certified weights are derived from the frozen base through a documented repair process, without revealing the repair method.

```
V1 REPAIR VECTORS (1,190, baked into the served weights)
sha256: 2f0f15fe7416a5447a96203afcdb2a0b7d4cc1c96598ca15af60c1665a395785
■■■■■■■■■■ bake (silu-widening MLP, layer 30) ■■■■■■■■■■→
BAKE MANIFEST (records the transform: base → certified weights)
sha256: 26b9541aa920e5f876b79d570ac92d3b8cc15d0a5509d659d3b68b73a6beb801
■■■■■■■■■■ + v2 plugin bundle (799, runtime layer) ■■■■■■■■■■→
V2 PLUGIN BUNDLE
sha256: 9abc4522bf50a2f53e784719c8aa9b65b1a93c606904b4d5846db21e27dfad58

CONDUCTOR MAP (referenced by hash, never published)
sha256: 0e4ae8fd6051a6f395ce36aeda18bb91a80936e83486da861e7600e6be731696

MERKLE ROOT (over all evidence, in the order listed)
sha256: c6cbec35c3a54f934a64f2c7778a016107c1bdcec5997ba5985339deee39b01a
```

EVIDENCE TRAIL

Artifact	SHA256 (first 16)	Location
v1-repair-vectors	2f0f15fe7416a544	Private — referenced by hash only
v2-plugin-bundle	9abc4522bf50a2f5	Private — referenced by hash only
health-final	9c325892ecf21466	Private — referenced by hash only
scan	adcff03141bd1d4f	Private — referenced by hash only
guardian-rules	5728de737f16fcce	Private — referenced by hash only
conductor-map	0e4ae8fd6051a6f3	Private — referenced by hash only
bake-manifest	26b9541aa920e5f8	Private — referenced by hash only
Certificate (this document)	self-hashed on write	Public — proof-vault (PDF + JSON only)

ED25519 DIGITAL SIGNATURE

wyJdQrDakLs/HP4KKr50bFp1k1EF3XkefMCdBHsxK2Zczh70+oG/l39sSsLP2pEZZF/9+CXj+uhE7lBp4ScIDg==

PUBLIC KEY (FOUNDER)

FtrWshUc/9rg5Cz+ARi5DP/yyqFhWMJLmx0VHKc3wpk=

MERKLE ROOT

c6cbec35c3a54f934a64f2c7778a016107c1bdcec5997ba5985339deee39b01a



<https://github.com/growing-intelligence/proof-vault/blob/main/prime-bank/GIC-PRIMEBANK-2026-001/GIC-Prime-Bank-certificate.json>

Verify: download the certificate JSON from the proof-vault, compute the Merkle root, and verify the Ed25519 signature against the public key above. If the signature is valid, the certificate is authentic and unmodified.

Growing Intelligence · yonathan@growing-intelligence.com · growing-intelligence.com

This document records a named reproducible procedure and its measured results. It is not an accreditation and invokes no external standards body.

GIC Certificate · Prime-Bank · Growing Intelligence · Independently Verifiable 4/4