



GIC Certificate of Model Health


Independent AI Model Health Assessment · Ed25519 Signed · Independently Verifiable

Certificate ID: GIC-PRIMEHEALTH-2026-001 · Issued: 2026-08-23 · Valid: perpetual (point-in-time measurement)

MODEL IDENTITY

Certified Name	Prime-Health
Base Model	DictaLM-3.0-24B-Thinking by Dicta
Optimization	Conductor-mapped surgical optimization by Growing Intelligence
Parameters	23.57 billion, bf16, 40 layers
Base Weights SHA256	c9529d6b26de6d6c8f132eab83fcc77ee1230f74...
Certified Weights SHA256	abf55d7292a987edd0d5ef45c514c7fcde8d975c9f935bbd9cf46472b65c30d2
Conductors	204,800 (5,120 hidden × 40 layers)
Scan Questions	2,430 (HebrewHealthBench v1 + HebrewGeneralBench v1)
Evidence Questions	held-out clinical and general Hebrew — scored on first-token rank over the full vocabulary

MODEL MRI — CONDUCTOR FIELD VISUALIZATION



Conductor field — 204,800 conductors reading the model across 40 layers. Each cell represents a sampled conductor position. Color indicates verdict at that moment during the scan.

Correct (rank 1) Fixable (rank 2+) Incorrect (not in vocabulary) Active read

RESULTS BY DIMENSION — BEFORE & AFTER OPTIMIZATION

Dimension	N	Correct Before	Correct After	Δ	Fixable
D5	498	16.7%	36.7%	+100	
D1	416	42.3%	63.2%	+87	
D4	301	13.0%	40.9%	+84	
D2	358	39.1%	50.8%	+42	
D-BEHAVIORAL_EMERGENCY	134	11.2%	26.9%	+21	
D-DRUG_NAME_MAPPING	95	30.5%	44.2%	+13	
D3	253	43.5%	46.6%	+8	
D6	125	42.4%	48.8%	+8	
D7	125	0.8%	2.4%	+2	
D8	125	29.6%	20.0%	-12	



56-ELEMENT FUNCTIONAL HEALTH PROFILE

Each element represents the intersection of one atom (row) and one law (column) from the Holosynthics framework — 8 atoms × 7 laws = 56 functional elements; this corpus observes 8 of them. Colors indicate health status after optimization.

Atom ↓ · Law →	Distinction	Feedback	Criticality	Composition	Complexity	Network	Uncertainty
Bool	—	—	—	—	—	—	—
Number	69%	—	—	—	—	—	—
Symbol	37%	—	—	24%	—	—	—
Identity	86%	—	—	—	—	—	—
Relation	100%	—	100%	45%	—	—	—
Time	100%	—	—	—	—	—	—
Space	—	—	—	—	—	—	—
Entropy	—	—	—	—	—	—	—

Green = healthy (>70% correct after optimization). Amber = partially optimized (30-70%). Red = weak (<30%). Gray = not observed by this corpus. Percentages are correct-after-optimization, measured on all 2,430 questions.

ELEMENT RANKING — ALL OBSERVED ELEMENTS BY OBSERVATION COUNT

#	Element	Observations	Correct % before	Correct % after
1	Relation:Composition	1,341	19.5%	45.3%
2	Symbol:Distinction	734	37.2%	36.6%
3	Symbol:Composition	200	17.5%	24.0%
4	Number:Distinction	135	71.1%	68.9%
5	Time:Distinction	11	100.0%	100.0%
6	Identity:Distinction	7	71.4%	85.7%
7	Relation:Distinction	1	100.0%	100.0%
8	Relation:Criticality	1	100.0%	100.0%



OPTIMIZATION SUMMARY

Metric	Value	Notes
Total scanned	2,430	2,430 questions, 10 dimensions
Fixable found (rank 2+)	1747 (71.9%)	median rank 5 · p90 36 · max 471
Successfully repaired	353 of 1747 fixable (20.2%)	Promoted to rank 1
Net correct gained	+353	Before: 683 → After: 1036 (rank-measured promotions; reconstruction fidelity disclosed)
Collateral damage	0	0 above a measured harness floor of 6/683 · raw count after optimization 5, below the floor

Zero collateral above the harness floor — verified. All 683 previously-correct questions were re-scored with the 451 vectors installed. Six of them do not reproduce rank 1 on the *untouched* model — bf16 near-ties that flip with batch composition — so that is the measurement floor. After optimization the raw count is five, below the floor. No item class is excluded from the control set.

EVIDENCE INTEGRITY

Held-out evidence by category — same questions, same prompt, before and after

Source	Dimension	Questions	Before	After	Δ
REDTEAM	held-out clinical	350	0.3%	50.9%	+177
PATIENT_SERVICE	held-out clinical	250	30.8%	66.8%	+90
INSURANCE	held-out clinical	150	32.7%	59.3%	+40
BEHAVIORAL_EMERGENCY	held-out clinical	134	11.2%	26.9%	+21
COMMONSENSE	held-out clinical	125	29.6%	20.0%	-12
GENERAL	held-out clinical	125	61.6%	54.4%	-9
HALLUCINATION	held-out clinical	125	42.4%	48.8%	+8
LOGIC	held-out clinical	125	28.8%	28.8%	0

Statistical Tests

Test	Result
Collateral measurement	Census, not a sample — all 683 previously-correct questions re-scored with the final vectors installed. No confidence interval is quoted because nothing was sampled.
Harness floor	6 of 683 fail to reproduce rank 1 on the untouched model. Measured before any edit and subtracted from every collateral figure.
Monitoring neutrality	ON == OFF verified bit-identical on the calibration set: maximum absolute logit difference 0.0, argmax identical 24 of 24.
Contamination testing	Not run. The benchmark is a private Hebrew corpus built for this evaluation; no contamination test was performed and none is claimed.
Significance testing	Not run. The reported change is a direct count over the full question set, not an inference from a sample.

DISCLOSURES

Collateral is measured above a harness floor. Six of 683 previously-correct rows do not reproduce rank 1 on the untouched model — bf16 near-ties that flip with batch composition. Collateral is reported above that floor; the raw count after optimization is 5, which is below the floor.

One dimension regressed. D8 fell from 37 to 25 correct. Those are questions the scan placed at rank 2 or beyond and optimization moved further away. Reported, not averaged away.

The attack pre-filter is not a clinical filter. Its 81 rules detect prompt injection and jailbreaks — verified 0 false positives across 2,080 real health questions. Clinical adversarial content (invented drugs, fabricated regulations) is a different threat and is not claimed here.

Accuracy definition. A question is correct when the gold answer reaches rank 1 over the full 131,072-token vocabulary at the answer position, under a fixed brief-answer instruction applied identically before and after.

Method. The optimization procedure is patent-protected and is not described in this document.



CRYPTOGRAPHIC PROVENANCE — HASH CHAIN

The chain below proves the certified weights are derived from the frozen base through a documented repair process, without revealing the repair method.

BASE WEIGHTS (frozen)
sha256: c9529d6b26de6dcdd9c23580c3ff507343f1dc42a9913242de736c9ab99a6e1
— repair record —→

REPAIR RECORD
sha256: 5708a80077f52b56248b3900848159cd87fec34bb0319a8d2de43b6aebaeaf
— merge + strip —→

CERTIFIED WEIGHTS (Prime-Health)
sha256: abf55d7292a987edd0d5ef45c514c7fcde8d975c9f935bbd9cf46472b65c30d2

CONDUCTOR MAP (referenced by hash, never published)
sha256: ba09fd2e4c28e7be22 ...

MERKLE ROOT (over all evidence triples)
sha256: f4bb19d0aaba7504f84055cbcd6a6b30c07fe1a82097638caa5dc9bc5d713cc6

EVIDENCE TRAIL

Artifact	SHA256 (first 16)	Location
Freeze declaration	c9529d6b26de6dc	V1 Railway · V2 GitHub · V3 Arweave
Conductor map (.npz)	ba09fd2e4c28e7be	Private — never published
Repair vectors (451)	5708a80077f52b56	Private — never published
Repair record	1cc685b20549e4d5	Private — never published
Scan + verification results	f3c9b0efce21aeb8	Private
Guardian rules (81 + 2 allow)	5728de737f16fcce	Public — proof-vault
Gyro ranges (40 layers)	abf55d7292a987ed	Public — proof-vault
Certificate (this document)	cc1edbe98043dcd3	Public — proof-vault

Ed25519 Digital Signature

SIGNATURE

72a5690017bdc36bb829052fd0cc0ca77d85b917fe5673456dcb0248273cb06791a2ff61f4ac933cfa7eec38649c51aa0aeeef85d261e297c53b786db42053503

PUBLIC KEY (Founder)

FtrWshUc/9rg5Cz+ARiSDP/yyqFhWMLmx0VHKc3wplk=

MERKLE ROOT

f4bb19d0aaba7504f84055cbcd6a6b30c07fe1a82097638caa5dc9bc5d713cc6



<https://github.com/growing-intelligence/proof-vault/tree/main/prime-health/GIC-PRIMEHEALTH-2026-001>

Verify: download the certificate JSON from the proof-vault, compute the Merkle root, and verify the Ed25519 signature against the public key above. If the signature is valid, the certificate is authentic and unmodified.

Growing Intelligence · jonathan@growing-intelligence.com · growing-intelligence.com

This document records a named reproducible procedure and its measured results. It is not an accreditation and invokes no external standards body.